

Estimating the Date of First Publication in a Large-Scale Digital Library

David Bamman
School of Information
University of California, Berkeley
dbamman@berkeley.edu

Michelle Carney
School of Information
University of California, Berkeley
michelle.carney@berkeley.edu

Jon Gillick
School of Information
University of California, Berkeley
jongillick@berkeley.edu

Cody Hennesy
Doe Library
University of California, Berkeley
chennesy@library.berkeley.edu

Vijitha Sridhar
Computer Science Division
University of California, Berkeley
vsridhar@berkeley.edu

ABSTRACT

One prerequisite for cultural analysis in large-scale digital libraries is an accurate estimate of the date of *composition* of the text—as distinct from the date of *publication* of an edition—for the works they contain. In this work, we present a manually annotated dataset of first dates of publication of three samples of books from the HathiTrust Digital Library (uniform random, uniform fiction, and stratified by decade), and empirically evaluate the disparity between these gold standard labels and several approximations used in practice (using the date of publication as provided in metadata, several deduplication methods, and automatically predicting the date of composition from the text of the book). We find that a simple heuristic of metadata-based deduplication works best in practice, and text-based composition dating is accurate enough to inform the analysis of “apparent time.”

CCS CONCEPTS

•Information systems → Digital libraries and archives;

KEYWORDS

Digital libraries; bibliographic metadata; publication date prediction

ACM Reference format:

David Bamman, Michelle Carney, Jon Gillick, Cody Hennesy, and Vijitha Sridhar. 2017. Estimating the Date of First Publication in a Large-Scale Digital Library. In *Proceedings of ACM/IEEE Joint Conference on Digital Libraries, Toronto, Canada, June 2017 (JCDL'17)*, 10 pages. DOI: 10.475/123.4

1 INTRODUCTION

The rise of large-scale digital libraries—such as those by Google Books, the Internet Archive, and the HathiTrust—has enabled a range of work in cultural analytics over the past decade, helping

provide the raw material for the historical analysis of genre [47], character [1], emotion [18], loudness [25], geographic attention [48] and much more.

For all of this work, it is important to have a rich understanding of a corpus prior to drawing conclusions about it; one important feature for understanding texts is their date of publication, since the primary variable in cultural analysis is often some quantity (such as word frequency) anchored specifically in time. For example, Michel et al. (2010) [32] was one of the first studies to make use of these extensive resources, and measure “fame” (among other quantities) by tracing the frequency of mention of a person’s name over the scope of their collection. Time is critically important for cultural analysis within these datasets, since arguments often hinge on exactly when a word was written, and criticisms may arise at the uncertainty of that information [37, 40]

For books, however, time can be measured in different ways. As the Functional Requirements for Bibliographic Records (FRBR) [20] model articulates, books can be viewed in several abstract categories, and each of those categories may have different temporal information associated with it—including the publication date of a specific edition (a *manifestation*, in FRBR terminology), and the first appearance of any expression of the *work* overall. The public domain texts of the HathiTrust, for example, include an edition of Jane Austen’s *Pride and Prejudice* published in 1920 by D. M. Dent & Sons as a reprint of an edition originally published in 1906; all of these texts are different FRBR expressions of Austen’s work, first published a century earlier in 1813. If we are using Austen’s text as a source for quantities in cultural history (such as the first appearance of a particular word), then we would likely associate the original date of publication of 1813 with the text, rather than the date of printing of any later editions. If, on the contrary, we are using Austen’s text to investigate questions of typography, then we may be more likely to associate the publication date of the edition. Different research questions require different associations of date.

This disparity between the observed date of publication for a specific printing and the (often unobserved) date of first publication is exacerbated for popular texts that undergo several reprintings. Figure 1 illustrates this for the works of Jane Austen in the public domain of the HathiTrust; nearly all of Austen’s works in this digital library are published after her death in 1817, most nearly a century later.




Figure 1: Distribution of publication dates for books in the HathiTrust for whom Jane Austen is the author. Austen’s death is marked at 1817.

One solution to this problem would fall out naturally from the process of deduplication: when deduplicating books, we simply retain the work with the earliest date of publication observed among its near-duplicates. However, as Nurmikko-Fuller et al. (2015) [38] note, this leaves open two problems: first, successful dating is dependent on the entire collection being complete (with the earliest published version of a book present in the collection); in the HathiTrust, for example, the earliest edition of *Pride and Prejudice* is the third edition (published in 1817); the earliest edition of Austen’s *Persuasion* is an edition from 1863 (46 years after its original publication). Second, deduplication itself is a heuristic that may also introduce errors of its own (either errors of omission in failing to recognize duplicates, or errors of commission in falsely attributing two unlike books to be the same).

In this work, we present an empirical evaluation of these methods. Our contributions are the following:

- We present a manually annotated, gold standard dataset for the dates of *first* publication for 2,706 books in the public domain of the HathiTrust, which we are publicly releasing under a Creative Commons license for others to use as a benchmark.
- We find approximately 10% of books of the HathiTrust have a difference of at least 10 years between their first publication date and observed publication dates in the metadata, suggesting the overall difference is large enough to warrant consideration.
- We find that when digital library metadata is comprehensive enough, simple deduplication-based methods are often accurate enough to substantially reduce this error.
- In addition to supplementing metadata-based assignment methods, estimating the date of publication from the *content* of a book can also yield insight into the determinants of temporal signatures in a text.

2 MANUAL DATING

To help drive an empirical evaluation of the accuracy of several methods for identifying the date of publication of a FRBR *work*, we create three gold standard datasets of books drawn from the HathiTrust paired with their manually identified date of first publication. From all books in the public domain of the HathiTrust, we draw three separate samples of approximately 1,100 books each.

- **UNIFORM.** We draw a uniform random sample of books from the public domain of the HathiTrust.
- **FICTION.** We draw a uniform random sample of books from the 102,349 works automatically identified to be fiction by Underwood 2014 [46].
- **STRATIFIED.** We draw a random sample of books from the public domain of the HathiTrust, but stratify the sample by observed publication date, so that we have a roughly equal number of books from each decade between 1750-1922.

We then attempt to manually identify the first publication date for each of these books using the process described below.

2.1 Process

Despite significant advances in implementing the FRBR model in bibliographic catalogs such as WorldCat [19], there is still no comprehensive, authoritative database that will merge or visualize a FRBR work with all of its known expressions and manifestations. To complicate matters, the bibliographic records that we depend on for documentation of publication dates were created according to standards and practices that have shifted over time [7], and have varied by institution [36].

In consideration of the above challenges, we consulted bibliographic records from a variety of sources, as well as digital books from a few major repositories, and assigned earlier dates of publication whenever there was reasonable documentary evidence to suggest that the year provided for a given book in our sample was not the earliest known edition. Any assigned date in our dataset should not be cited as an authoritative judgement on that book’s earliest publication, but rather as an indication that the bibliographic record suggests there was an earlier publication date than the edition at hand. The sources we consulted in the process were:

- Bibliographic records from the *Nineteenth-Century Short Title Catalogue (NSTC)*, available in C19: The Nineteenth Century Index (ProQuest).
- Library catalog records from WorldCat (OCLC).
- Date and edition notes in the front matter and prefatory remarks from other digital editions of the title available in HathiTrust, Google Books, and Eighteenth Century Collections Online (Gale).
- Less frequently, and for more difficult titles, we referred to encyclopedias, descriptive bibliographies of specific authors, and/or book reviews from periodicals indexed in C19.

The *NSTC* unites records for English language printed works published between 1801 and 1918 from eight major libraries where, as noted in the introduction to the first series, cataloging standards varied widely [36]. While enumerative bibliographies such as the *NSTC* provide an imperfect or incomplete record of a book’s publication history [10, 11], book historians note their importance and relative accuracy for large-scale bibliographic scholarship [12, 45]. The majority of the bibliographic records in the *NSTC* were created after “distinguishing different editions of the same work had become the norm in bibliographic control” and before a shift in cataloging standards towards FRBR and a renewed focus on the work [7]. Most publication dates in the *NSTC* were accordingly

drawn directly from a book's title pages. Records for some undated imprints, however, appear to have used the nearest approximate beginning or middle year of a decade (e.g., 1850 or 1855) as an approximate publication date without documenting that fact [12].

Bibliographic records in WorldCat offer two distinct advantages to those in the *NSTC*: first, modern cataloging guidelines state to clearly document when dates used were approximate, based on any sources outside of the book itself, or represented dates other than that of the imprint (e.g., copyright) [15]. Second, WorldCat contains over 220 million records for books, compared to 1,278,000 in C19's *NSTC*, and is often a better source for obscure titles. WorldCat also allows for the compilation and display of the many FRBR expressions of a given work by publication date, providing a quick, if incomplete, view of a book's publication history. *NSTC*, however, remained the better source for dating popular works: it is far easier to locate the 1813 first edition of *Pride and Prejudice* among the 80 odd records in *NSTC*, for example, than it is from more than 4,000 records of expressions of the title in WorldCat.

2.2 Design choices

By grounding books in large-scale digital libraries in their earliest date of publication, we hope to enable more precise measurements of cultural phenomena that are realized in the language of the text; this fundamental use case drives several design choices in our manual dating process.

According to the FRBR model, the record for an expression of a given work (a translation, for example) would note both the publication date of a work in its original language as well as the date of the expression being cataloged. Since this has not been a consistent bibliographic practice, however, and because it goes beyond the scope of our interests here, we have chosen to annotate a single date for each item in our sample, following the guidelines below.

Translations. We have assigned translations that are the work of named translators in our sample with the earliest known date of the translation's publication, not the publication date of the work in its original language. *The Attic Nights*, originally composed by Aulus Gellius sometime before his death in 180 AD, and translated into English by W. Beloe in 1795, is assigned the date of 1795, for example.

Serials. Portions of novels, as well as individual poems, short stories and essays, are often published in magazines and newspapers for years before their first publication as a unified work. We have dated the first printed editions of complete works rather than any individually printed components whenever there is clear evidence that these are author-initiated publications intended to appear as complete works.

Edited material. Popular works are often published and republished with prefatory materials, notes, and back matter composed at various times, and by a variety of authors and editors. We have chosen to select the date that best reflects the first publication of the **core work**, ignoring the presence of supplementary text.

2.3 Exclusions

There were several categories of books that could not be assigned any meaningful date of first publication, and so were not manually

dated and are thereby excluded from our sample. A substantial number of volumes defy the notion of encompassing a single work (compilations, for example). Others were poor candidates for large-scale text analysis, since they included too few textual elements or were in formats leading to poor OCR.

Compilations. Compilations may contain works composed or first published across a range of dates in non-serial publications, and therefore have not been dated in our sample. Included in this category are primary sources later collected and published: *The Letters of Abelard and Eloisa* (letters), *An Account of the Life of the Reverend David Brainerd* (diary entries), and *Speeches, Messages, and Other Writings of the Hon. Albert G. Brown* (speeches). Essays that are compiled by an editor or publisher are also included (e.g., *Miscellaneous Essays, Volume Two of Modern British Essayists*, by Archibald Alison). By far the most representative are collected works previously published in non-serial publications: e.g., *The Dramatic Works of Beaumont and Fletcher* (plays), Anton Chekhov's *The Kiss, and Other Stories* (short stories), or Riverside Press' 1883 *The Complete Works of Nathaniel Hawthorne, Volume 3* (novels). Poems previously published over a wide range of dates (e.g., *The Poetical Works of S.T. Coleridge*) fall in this category while those released as a single volume by an author (e.g., *Poems in Oil and Other Verse*, by Will Ferrell), and book-length poems (e.g., *The General* by Francis Gentleman) do not. This category does not include collections of short stories for which we were not able to find clear evidence of their being previously published in non-serial titles (e.g., *Stories of New England Life; or Leaves from the Tree Igdrasil* by Martha Russell). Compilations were found to be quite pervasive in all of our samples, accounting for approximately 20% of all books we attempted to date.

Abridgments. These texts are revised versions of another original work, often including much of the same text verbatim. We have chosen not to date these works since it is unclear how much of the original source is included. *The Student's Hume: A History of England* by David Hume (1859), for example, "[incorporates] the corrections and researches of recent historians," but those corrections are undocumented in the text.

Folk tales. This category includes common stories for which there is no clear attribution of an individual author. George Routledge and Sons undated edition of *Cinderella*, for example, is only one of thousands of expressions of this tale.

Non-text. This category includes sheet music (e.g., *Troubadour Song* by Harriet Browne), tables (e.g., *Observations on the State of the Air, Winds, Weather* by Joseph Dymond), directories, lists, maps, dictionaries (e.g., *Chinook Vocabulary, Chinook-English*), images, manuscript facsimiles (e.g., *The World and the Child*), or non-English languages. Approximately 5% of books in our collection were found to belong to this category.

3 ANALYSIS

After excluding books for dating using the criteria above, the final annotated dataset contains 916 manually dated books in the UNIFORM sample, 991 books in the FICTION sample, and 799 books in the STRATIFIED sample. For the experiments reported below,

we hold out 500 books from each sample for evaluation and use the remainder for training and model selection (including hyperparameter optimization).

| Sample | Train | Test |
|------------|-------|------|
| Uniform | 416 | 500 |
| Fiction | 491 | 500 |
| Stratified | 299 | 500 |

Table 1: Size of the different samples.

How much of a difference is there between the true dates of first publication and the publication dates of specific printings? Table 2 presents these results for all of the labeled data, along with 95% confidence intervals calculated using the bootstrap [9]: on average, the absolute difference among books in the uniform dataset is 2.73 years; among fiction books, the difference is approximately 4.66 years, and among those in our sample stratified by decade, the difference is approximately 5.09 years. While the mean combines the large majority of books with no (or minimal) distance along with outliers with an extremely large difference (e.g., several hundred years), we find that approximately 10% of books across our samples have a distance greater than 10 years; 1% of them have a difference greater than 100 years.

| Sample | Mean absolute Δ | ≥ 10 | ≥ 25 | ≥ 100 |
|------------|------------------------|-----------|-----------|------------|
| Uniform | 2.73 [2.1, 3.48] | 7.1% | 3.3% | 0.3% |
| Fiction | 4.66 [3.49, 6.12] | 10.3% | 5.5% | 0.8% |
| Stratified | 5.09 [4.01, 6.53] | 10.8% | 5.3% | 0.9% |

Table 2: Difference between the manually identified first publication date and that reported in the metadata, as measured by mean absolute error (with 95% confidence intervals), and percent of books with a difference of at least 10 years, 25 years and 100 years.

4 AUTOMATIC DATING

4.1 Metadata

The metadata provided in MARC records that attend books in the HathiTrust can be used to estimate the date of first publication in several ways. Date information in MARC records generally appears in the 008 subfield (appearing as volume-specific *Date 1* and *Date 2* regions (whose semantics depend on the type of publication they appear with) and can also appear in custom z30 tags. The HathiTrust determines the publication date as part of a rights-determination algorithm.¹ For reprints, the earlier of the two dates (captured in *Date 1*) is used; for works whose publication date is also the copyright date, that date is used; for all other works, the later of all observed dates (in any field) is used. Since this algorithm prioritizes the determination of rights in assigning dates, it is generally conservative, preferring the later dates observed. We term this below as HATHI DATE.

¹Detailed here: https://www.hathitrust.org/bib_rights_determination.

Without using dates for the determination of rights, we can adopt an alternative method for estimating the date of first publication by prioritizing the earliest dates mentioned in the 008 or z30 fields. For non-series, we simply take the earliest date mentioned among all dates observed; for series, which often list the date of publication of the first volume in the series as *Date 1*, we take as our prediction the earliest date mentioned in the z30 field. This has the effect of biasing our estimate earlier. We term this below as EARLY.

4.2 Metadata-based deduplication

Outside of the individual metadata records for books in the collection, we can also leverage the resource of the collection as a whole. The HathiTrust, like many large-scale digital libraries, contains works aggregated from several university libraries [23]; each of these libraries may contain multiple printings and editions of a given expression for a work, and we can leverage this pattern of duplication to find better estimates of the first date of publication.

To do so, we define a metadata-based duplicate with a simple and reproducible heuristic: two metadata records are duplicates if the first 25 characters in both the title and author fields are the same. The restriction to the first 25 characters mitigates some variation as a function of the granularity of titles (e.g., *The life and adventures of Robinson Crusoe, of York, mariner* vs. *The life and adventures of Robinson Crusoe*). For any given work, we assign its date of publication to be the earliest among its identified duplicates. We term this below as METADATA DEDUP.

4.3 Content-based deduplication

Metadata-based similarity can fail in several ways. First, requirements on the similarity of titles and authors may not identify as duplicates those works with slight variants (e.g., *The life and adventures of Robinson Crusoe* vs. *Robinson Crusoe*). Second, it may not identify at all works whose metadata is erroneous, or whose works are found in collections or serials (e.g., *Great Expectations* vs. vol. 7 of Charles Dickens’ *Works*).

Content-based deduplication provides an alternative that has been used extensively in digital libraries [42–44, 49, 50]. To identify duplicates based on the content, we represent books as a set of *shingles* comprised of hashed word trigrams, leveraging minhashing [2] to generate a compact 500-dimensional representation. We then identify near-duplicates in the entire 5.6 million book dataset using locality-sensitive hashing [21] on those minhash representations (placing each book into 250 distinct buckets, each identified using two sequential minhash features). Near-duplicates are defined as those with a jaccard similarity above some threshold k , and we set $k = 150$ using cross-validation on the training partition of the gold data described in section 2. We term this below as CONTENT DEDUP.

A third deduplication-based approach is to leverage both metadata deduplication and content deduplication in a single model, assigning the earliest date generated by either process. We term this COMBINED DEDUP below.

| Sample | Hathi Date | Early | Metadata dedup | Content dedup | Combined dedup |
|------------|------------|-------|----------------|---------------|----------------|
| Uniform | 3.12 | 2.88 | 2.36 | 2.60 | 2.74 |
| Fiction | 4.68 | 4.44 | 3.12 | 3.52 | 2.82 |
| Stratified | 5.01 | 4.63 | 3.46 | 4.08 | 3.39 |

Table 3: Mean absolute error for predicting the date of first publication in the test partition of the gold data.

4.4 Results

Table 3 illustrates a comparison of the metadata- and deduplication-based methods. While the EARLY methodology on its own uniformly improves upon the dates of publication provided by the HathiTrust by their rights-determination algorithm for this particular purpose, we see that deduplication methods yield substantial improvement, leading to a 25-30% reduction in error over metadata alone. While COMBINED DEDUPLICATION performs the best on two of the samples, the simpler method of deduplicating based on the metadata alone performs best for the UNIFORM sample and is competitive across all three samples.

5 CONTENT-BASED PREDICTION

Metadata and deduplication-based methods of establishing the date of first publication depend on having either high-quality metadata or relatively large collections in which meaningful duplicates can be found. To explore the possibility of predicting the first date of publication from the content of the text itself, we assess two different approaches: one based on existing aggregate information from Google Books (in the context of a Naive Bayes classifier), and one based on training a discriminative model (linear regression with ℓ_2 regularization) directly on a large sample of books from the HathiTrust.

For both models, we represent a book as a bag of unigrams, but only consider text that appears after the first 10% of pages at the start of the book (and a minimum of 10 pages), and before the final 6% of pages at the end of the book (which, in our data, were the average beginning and end of the core content of a book). This encourages the models to only have access to linguistic features of the text, and not any bibliographic data present in the frontmatter or general paratext (such as advertisements or publisher’s information at the end). This design choice allows us to assess the degree to which content-based prediction is complementary to metadata-based prediction (and not using similar information in making predictions).

5.1 Google Ngrams + Naive Bayes

One method for estimating a likely date of publication is to use existing aggregate resources from large-scale digital libraries, such as Google Books. The Google Ngram dataset contains ngram counts (unigrams up to 5-grams) for a total of 4.5 million books published between 1505 and 2008, in effect providing us with pre-calculated language models for each year in their collection (though the year of publication is that for specific editions, and not the date of first publication of any expression of the work).

With a Naive Bayes model, we make a categorical prediction for a year $y \in \mathcal{Y} = \{1700, \dots, 1922\}$ given a set of observed words in a

book W , our estimate of the prior belief over publication dates absent any data (θ) and a set of unigram language models ϕ , one for each year in the label space $\{\phi_i \mid \forall i \in 1700, \dots, 1922\}$. Each ϕ_i lies in the V -dimensional probability simplex; in the experiments that follow, $V = 100,000$ (the 100,000 most frequent words in the data overall). Since Google Ngrams provides count data for terms by year, we simply set ϕ to be the empirical frequency of those terms as observed in that year, plus a small amount of additive smoothing, normalized by the total counts for that vector. Since our prior belief about the likely years differs according to our sample (e.g., our uniform sample and decade-stratified sample have different distributions of years by design), we estimate it empirically from the years observed in the training data.

This approach presents two issues: first, we are treating y as a categorical variable rather than the ordinal one we know it to be. Second, the Google Ngram data is quite variable when aggregated at the level of any individual year, and reflects idiosyncrasies in the samples of which books are published (and how many of them). As figure 2 illustrates, the raw relative frequency of the term *thee* has substantial fluctuation, especially among the less-well-attested years before 1800. This variability is such that the term is used twice as frequently in 1717 than in either 1716 or 1718.

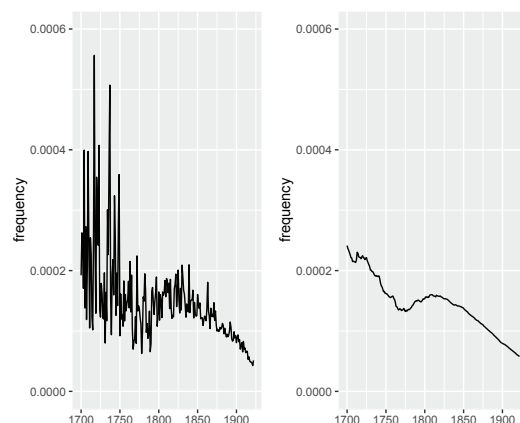


Figure 2: Distribution of the relative frequency of *thee* over time, raw (left) and smoothed with a 50-year moving average (right).

Additive smoothing is one way of accounting for this variability, corresponding to adding some small mass (either constant in the case of Laplace smoothing, or variable, in more general Bayesian case of drawing a multinomial from an asymmetric Dirichlet). In our case, however, we have strong reason to believe

| Sample | Baseline | NB | Conv NB | Ridge reg. |
|------------|----------|-------------------|-------------------|------------------|
| Uniform | 22.9 | 16.3 [14.9, 17.8] | 17.8 [16.2, 19.6] | 10.1 [9.4,10.9] |
| Fiction | 25.7 | 18.0 [15.8, 20.6] | 15.9 [14.1, 18.2] | 8.1 [6.4,10.6] |
| Stratified | 46.3 | 29.2 [26.8, 31.8] | 23.6 [21.4, 25.8] | 24.5 [22.7,26.4] |

Table 4: Mean absolute error for predicting the date of first publication in the test partition of the gold data.

that the unigram language models of adjacent time periods should also be similar (e.g., $\phi_{t-1} \approx \phi_t \approx \phi_{t+1}$). To encode this assumption into our model, we convolve our empirical estimates of ϕ with a K -dimensional kernel k , equivalent to calculating a moving average over ϕ along the time dimension, producing a new set of temporally smoothed unigram language models ϕ^* .

$$k = \left[\frac{1}{K} \right]^K$$

$$\phi^* = \phi * k$$

We then calculate the probability of year y for a book with words w through the usual application of Bayes’ rule:

$$P(y \mid \theta, \phi^*, w) \propto P(y \mid \theta) \prod_w P(w \mid \phi_y^*)$$

In this model, K (the convolution size) is a tunable parameter, and we optimize it on the training data.

5.2 HathiTrust + Ridge regression

Naive Bayes makes a strong simplifying assumption: that all features are independent and contribute equally to the prediction. This can be problematic in some scenarios where a large number of features are highly correlated with each other, leading to artificially high confidence in a prediction simply due to the fact of repeated but not particularly informative features.

In using an existing dataset, we may also be hampered by the corpus selection process; without knowing exactly how the Google Ngram counts were generated (and, in particular, which books they were generated from), we lose some control over our analysis. (In a very real possibility, the ngram counts could be generated from non-deduplicated data, where many editions of the same work or even copies of the same edition in different libraries, are treated as different books in generating the counts.)

To account for both of these factors, we also train a discriminative linear regression model with ℓ_2 regularization (ridge regression) using books from the HathiTrust.

We train using three different datasets, each corresponding to one sampling strategy described above. The UNIFORM dataset is a sample of 100,000 books selected at random from all the works in the public domain of the HathiTrust, but—importantly—excluding all books written by the same author as any book in our test dataset (which thereby also excludes all of our test books from being included in the training sample). We treat all authors who share the same first 25 characters of their name as referring to the same individual. This exclusion yields a total dataset of 43,517 books.

The FICTION dataset includes all books in the Underwood 2014 [46] fiction dataset, with the same exclusion criteria as described above. This yields a total of 54,802 books.

The STRATIFIED dataset includes a random sample of all books in the HathiTrust, but stratified by their date of publication. Using the same exclusion criteria yields a total of 14,547 books.

For all books in the training set, rather than using the observed date of publication as the true label whose prediction we are trying to optimize, we use the metadata deduplication strategy discussed in §4.2 to assign its label to be the date of publication for the earliest work among its identified duplicates. We featurize each book as a set of binary indicators for all of words it contains that are also in the vocabulary of the 100,000 most frequent words overall, and again exclude all information that appears in the first 10% of last 6% of pages in the book.

5.3 Results

To compare the performance of the different models, we calculate the mean absolute error between the predictions they make on the test data \hat{y} and the manually identified gold labels for those books y :

$$\text{MAE} = \frac{1}{N} \sum_{i=1}^N |\hat{y}_i - y_i|$$

In order to quantify our uncertainty around this measure given our sample, we calculate 95% confidence intervals using the non-parametric bootstrap [9], using 1000 bootstrap resamples of the test data and calculating the mean absolute error for each.

Table 4 presents the results of this analysis. In order to help contextualize the results, a simple baseline that predicts the average value of all books in the training partition of the data yields a mean absolute error of 22.9 years for the UNIFORM dataset, 25.7 years for the FICTION dataset, and 46.3 years for the STRATIFIED dataset. While ridge regression and Convolutional Naive Bayes are equivalent for the harder problem of predicting the date of publication for the STRATIFIED dataset (for which there is both less training data and greater variance among the dates), ridge regression with binary features strongly outperforms the Naive Bayes models built from existing Google Ngram data for both the UNIFORM and FICTION samples, leading to a substantial reduction in error between 40-50%. Training a discriminative regressor directly on the full text of individual books, rather than relying on aggregate counts for entire years, leads to measurably stronger predictions.

6 ANALYSIS

As the automatic dating task above illustrates, the best method for estimating the date of first publication for books in a large digital library is to leverage the depth of the collection, identifying duplicates and assigning the first date of publication for a book to be the earliest date attested among its near-duplicates. While

| Author | Title | Narrative time | Date of publication | Predicted date |
|------------------------|----------------------------------|----------------|---------------------|----------------|
| Arnold Bennett | The Old Wives' Tale | 1872 | 1908 | 1914 |
| John Galsworthy | The Man of Property | 1870 | 1906 | 1907 |
| Winston Churchill | The Crossing | 1774 | 1904 | 1904 |
| Stephen Crane | The Red Badge of Courage | 1863 | 1895 | 1897 |
| George Moore | Esther Waters | 1875 | 1894 | 1897 |
| Robert Louis Stevenson | The Master of Ballantrae | 1745 | 1889 | 1878 |
| Marcus Clarke | For the Term of His Natural Life | 1827 | 1874 | 1879 |
| Elizabeth Gaskell | Sylvia's Lovers | 1790 | 1863 | 1865 |
| Charles Dickens | A Tale of Two Cities | 1794 | 1859 | 1867 |
| Walter Scott | The Bride of Lammermoor | 1707 | 1830 | 1851 |
| James Fenimore Cooper | Last of the Mohicans | 1757 | 1826 | 1836 |
| Walter Scott | The Heart of Midlothian | 1736 | 1818 | 1840 |
| Walter Scott | Rob Roy | 1715 | 1817 | 1834 |

Table 5: A sample of historical novels, along with their date of first publication, narrative time within the story, and predicted date according to our model.

automatically predicting the date of first publication from the *content* of a book is not as precise, it can still serve several important ends. As Guo et al. (2015) [16] point out, a substantial fraction of books in the HathiTrust (13%) are missing publication date metadata, and even when present, information in legacy metadata can often be improved [13]; automatically estimating the date of publication, or an interval in which publication is likely, can help with the search and discovery of these texts when date is provided as a facet.

Another end is to analyze the degree to which the *predicted* date of composition agrees or disagrees with a given date of publication. Practically, books whose predicted and given dates strongly disagree may be good candidates for quality assurance. They can also provide the raw material for an analysis of what a content-based date prediction system is actually learning. When making predictions, are we relying more on historical markers in the text (e.g., mentions of historical figures like *George Washington* or explicit dates like *1776* that anchor a text in the late 18th century) or are we relying more on linguistic and stylistic signals characteristic of the time in which the book was written (such as *thee* and *thou*)?

To illustrate this, we consider a small selection of historical novels from the FICTION dataset, where there is a significant difference between the date of composition and the narrative time within the book; for example, James Fenimore Cooper’s *Last of the Mohicans* was written in 1826 but narrates a story that takes place in 1757. To identify historical novels, we leverage a list of *Best Historical Fiction* on GoodReads² and identify the set intersection between that list and the books in our FICTION dataset, also adding the further constraints that the books must be written in English and describe a narrative time after 1700. Table 5 lists the 13 books that meet these criteria; this is small sample, and can best serve as an anecdotal case study of the potential for these methods.

To compare the predicted and given dates of publication for these historical novels, we train our best-performing model from table 4 above (ridge regression), in a ten-fold cross-validation of the full fiction dataset, taking care that works by the same author

never show up across folds (so that we do not train on one copy of *Great Expectations* in one fold and use that information to estimate the date for another copy in a different fold).

In doing so, we use a model trained on 90% of the data to make predictions for the held-out 10%, and iterate through ten folds to make predictions for the entire dataset. As table 5 shows, we see that the predictions our models makes for these historical novels much more closely align with their actual date of first publication than the imagined historical time within each novel; the markers of linguistic style rather than the historical content appear to be driving the prediction decision.

Given the design of our model, we can also ask what the most indicative terms are within each book that lead to the estimate. In using linear regression, we make predictions for a book, represented as a V -dimensional feature vector $x \in \mathbb{R}^V$, by calculating the dot product with a corresponding set of feature coefficients $\beta \in \mathbb{R}^V$:

$$y = \sum_{i=1}^V x_i \beta_i$$

Martens and Provost (2014) [31] provide one method for explaining binary *classification* decisions in linear models by asking the following question: what is the minimal set of features that x has that, if removed, would lead us to predict the opposite label? We can extend this line of reasoning to the case of regression as well (where we predict a continuous variable rather than a discrete category) by asking: given a prediction \hat{y} for a given data point x , what is the minimal set of features of x that we can remove in order to predict an alternative target \hat{y}' ? To illustrate with one simple case study, what are the set of features we can remove from *Last of the Mohicans* to change our predicted date from 1835 to the narrative time of 1757?

As in Martens and Provost (2014), we can do so simply by ranking the learned coefficients β and removing features from x with the strongest positive weights (to push the prediction earlier in time) or the strongest negative weights (to push the prediction later in time). Removing 328 of *Last of the Mohican*’s 8976 features (3.7%) yields a target prediction \hat{y}' of 1757. The ten strongest

²https://www.goodreads.com/list/show/15.Best_Historical_Fiction

weighted features are shown in table 6 and reveal the set of ahistorical terms that Cooper uses in the novel that are not only uncharacteristic of language spoken in 1757, but are also even on the leading edge of wider adoption, as we can confirm with an external data source in Google Ngrams. Figure 3 illustrates the inflection point between the use of *every one* and *everyone*, and figure 4 illustrates the precipitous rise of *later*, both of which Cooper is an early adopter (which also in part leads to a higher predicted date of publication of 1835 than its true date of first publication of 1826).

| Term |
|------------|
| etc. |
| later |
| everything |
| everybody |
| anything |
| ahead |
| lack |
| big |
| simply |
| meantime |

Table 6: Ahistorical terms in James Fenimore Cooper’s *Last of the Mohicans*.

Even though our estimates of the true first date of publication are better served with deduplication-based methods, learning a model to predict this date from the content of the book gives us the potential for deeper insight into the books in our collection by providing a mechanism for measuring *apparent time*, as distinct from both the observed publication date or the narrative time.




Figure 3: Relative frequency of *everyone* and *every one* in the Google Ngram data.




Figure 4: Relative frequency of *later* in the Google Ngram data.

7 RELATED WORK

Our approach and general methodology here are motivated and informed by much prior work in automatically predicting the date of publication for text, including newspaper articles [41], web pages [24, 39], Wikipedia biographies [28], literature (such as short stories from Project Gutenberg [29] and Romanian novels) [4], and large digital libraries [16, 30].

Nearly all of these approaches choose an experimental design used by de Jong et al. (2005) [8], which partitions time into discrete buckets (e.g., 1870-1880) and attempts to classify a new document into one of those buckets—rather than, as in our case, attempting to pinpoint the exact date of publication (as a regression problem). There are tradeoffs to both design choices: classification into discrete buckets allows for the natural incorporation of features based on the distribution of a language model expressed in that bucket [8, 16, 24, 29] at the cost of the specificity of the prediction (a model may be able to predict with finer detail than the bucket allows); regression is more limited in the information that can be brought to bear on the problem but potentially allows for finer-grained predictions, while also enabling the analysis discussed in section 6: understanding what features contribute to a prediction, and isolating the specific characteristics of a text that give its date away.

While other work has also focused on the prediction of publication dates in large digital libraries (either the Hathi-Trust [16] or Google Books [30]), one further contribution we make is distinguishing between the different meanings of “publication date” when seen in the context of the FRBR hierarchy: all work to date has leveraged the publication date provided in the metadata for training and evaluation, effectively focusing on learning the temporal characteristics at the FRBR level of a *manifestation*—when a specific edition was published. This variety of temporal metadata certainly has its use cases, but is only one choice among several: in creating a dataset and evaluating our predictions on the *first* date of publication, we are effectively assessing our ability to learn the temporal characteristics of books at the FRBR level of a *work*.

Related to work in the explicit prediction of the publication dates of text is adjacent work in estimating the change in the meanings of words over time [14, 17, 22, 26, 27, 33, 35]; while the goal of this work is often to characterize the linguistic dynamics of change, and to estimate when several different words across time all refer to the same (or similar) concept, it can also be used in the design of features that provide information not simply on the presence, absence or frequency of a specific lexical form (as in this and other past work), but also when a given lexical form is used in a way that is particularly characteristic of specific time periods (e.g., while *apple* is a common word throughout the history of print, it only recently is used in the context of a specific organization).

In characterizing the temporal signature of time periods and providing a way to measure the degree to which an author or a specific book predates the wider adoption of some linguistic phenomenon, one potential future direction for this work is in analyzing the diffusion of stylistic innovation, especially in the creating and popularization of neologisms [3, 5, 6, 34].

8 CONCLUSION

We present in this work a new gold standard dataset and several empirical analyses for predicting the *first* date of publication for books in a large-scale digital library, making an important distinction between the date a specific FRBR *manifestation* was published and the date its original *work* (of which it is an instantiation) was published.

As more and more work in the digital humanities, computational social science and cultural analytics is increasingly making use of the texts and their attendant metadata in large-scale digital libraries, this is an important distinction to be made; both measures are appropriate for different analyses, and, as we show, a simple metadata-based deduplication method is often acceptable for estimating the work publication from a set of manifestation dates in a large enough collection. When metadata is lacking, content-based estimation can also yield relatively accurate measures, and can itself occasion analyses on the linguistic and stylistic characteristics of books and the authors who write them.

All annotated data is available for public use under a Creative Commons license at: <https://github.com/dbamman/jcdl2017>.

9 ACKNOWLEDGMENTS

We thank the anonymous reviewers and Ted Underwood for valuable advice, and are grateful to the HathiTrust Research Center for their assistance in enabling this work. The research reported in this article was supported by a grant from the Digital Humanities at Berkeley initiative.

REFERENCES

- [1] David Bamman, Ted Underwood, and Noah A. Smith. A Bayesian mixed effects model of literary character. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics*, 2014.
- [2] A. Broder. On the resemblance and containment of documents. In *Proceedings of the Compression and Complexity of Sequences*, SEQUENCES '97, 1997.
- [3] Paula Chesley. *Linguistic, cognitive, and social constraints on lexical entrenchment*. PhD thesis, University of Minnesota, 2011.
- [4] Alina Maria Ciobanu, Liviu P. Dinu, Octavia-Maria Sulea, Anca Dinu, and Vlad Niculae. Temporal text classification for romanian novels set in the past. In *RANLP*, pages 136–140, 2013.
- [5] Paul Cook. *Exploiting linguistic knowledge to infer properties of neologisms*. PhD thesis, University of Toronto, 2010.
- [6] Chiru Costin-Gabriel and Traian Eugen Rebedea. Archaisms and neologisms identification in texts. In *RoEduNet Conference 13th Edition: Networking in Education and Research Joint Event RENAM 8th Conference, 2014*, pages 1–6. IEEE, 2014.
- [7] Karen Coyle. *FRBR, Before and After: A Look at Our Bibliographic Models*. ALA Editions, 2015.
- [8] F.M.G. de Jong, H. Rode, and D. Hiemstra. Temporal language models for the disclosure of historical text. In *AHC 2005*, pages 161–168, Amsterdam, The Netherlands, September 2005.
- [9] B Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, pages 1–26, 1979.
- [10] Simon Eliot. “Patterns and Trends” and the “NSTC”: Some initial observations. Part one. *Publishing History*, 42:79–104, 1997.
- [11] Simon Eliot. “Patterns and Trends” and the “NSTC”: Some initial observations. Part two. *Publishing History*, 43:71–112, 1998.
- [12] Simon Eliot. Very necessary but not quite sufficient: A personal view of quantitative analysis in book history. *Book History*, 5:283–293, 2002.
- [13] Katrina Fenlon, Colleen Fallaw, Timothy Cole, and M.J. Han. A preliminary evaluation of HathiTrust metadata: Assessing the sufficiency of legacy records. In *Digital Libraries*, 2014.
- [14] Lea Frermann and Mirella Lapata. A Bayesian model of diachronic meaning change. *Transactions of the Association for Computational Linguistics*, 4:31–45, 2016.
- [15] Michael Gorman. *The Concise AACR2*. American Library Association, 4th edition, 2004.
- [16] Siyuan Guo, Trevor Edelblute, Bin Dai, Miao Chen, and Xiaozhong Liu. Toward enhanced metadata quality of large-scale digital libraries: Estimating volume time range. In *iSchools Conference*, 2015.
- [17] William L. Hamilton, Jure Leskovec, and Dan Jurafsky. Diachronic Word Embeddings Reveal Statistical Laws of Semantic Change. In *Association for Computational Linguistics (ACL)*, 2016.
- [18] Ryan Heuser, Franco Moretti, and Erik Steiner. The emotions of london. Technical report, Stanford Literary Lab, 2016.
- [19] Thomas B. Hickey and Edward T. O'Neill. FRBRizing OCLC's WorldCat. *Cataloging & Classification Quarterly*, 39(3-4):239–251, 2005.
- [20] IFLA. *Functional Requirements for Bibliographic Records Final Report*. K.G. Saur Verlag, Munich, 1998.
- [21] Piotr Indyk and Rajeev Motwani. Approximate nearest neighbors: Towards removing the curse of dimensionality. In *Proceedings of the Thirtieth Annual ACM Symposium on Theory of Computing*, STOC '98, pages 604–613, New York, NY, USA, 1998. ACM.
- [22] Adam Jatowt and Kevin Duh. A framework for analyzing semantic change of words across time. In *Proceedings of the 14th ACM/IEEE-CS Joint Conference on Digital Libraries*, JCDL '14, pages 229–238, Piscataway, NJ, USA, 2014. IEEE Press.
- [23] Jacob Jett, Terhi Nurmikko-Fuller, Timothy W. Cole, Kevin R. Page, and J. Stephen Downie. Enhancing scholarly use of digital libraries: A comparative survey and review of bibliographic metadata ontologies. In *Proceedings of the 16th ACM/IEEE-CS on Joint Conference on Digital Libraries*, JCDL '16, pages 35–44, New York, NY, USA, 2016. ACM.
- [24] Nattiya Kanhabua and Kjetil Nørvåg. *Using Temporal Language Models for Document Dating*, pages 738–741. Springer Berlin Heidelberg, Berlin, Heidelberg, 2009.
- [25] Holst Katsma. Loudness in the novel. Technical report, Stanford Literary Lab, 2014.
- [26] Yoon Kim, Yi-I Chiu, Kentaro Hanaki, Darshan Hegde, and Slav Petrov. Temporal analysis of language through neural language models. In *Proceedings of the ACL 2014 Workshop on Language Technologies and Computational Social Science*, pages 61–65, June 2014.
- [27] Vivek Kulkarni, Rami Al-Rfou, Bryan Perozzi, and Steven Skiena. Statistically significant detection of linguistic change. In *Proceedings of the 24th International Conference on World Wide Web*, pages 625–635. ACM, 2015.
- [28] Abhimanu Kumar, Jason Baldridge, Matthew Lease, and Joydeep Ghosh. Dating texts without explicit temporal cues. *CoRR*, abs/1211.2290, 2012.
- [29] Abhimanu Kumar, Matthew Lease, and Jason Baldridge. Supervised language modeling for temporal resolution of texts. In *Proceedings of the 20th ACM International Conference on Information and Knowledge Management*, CIKM '11, 2011.
- [30] Yuanpeng Li, Dmitriy Genzel, Yasuhisa Fujii, and Ashok C. Popat. Publication date estimation for printed historical documents using convolutional neural networks. In *Proceedings of the 3rd International Workshop on Historical Document Imaging and Processing*, HIP '15, 2015.
- [31] David Martens and Foster Provost. Explaining data-driven document classifications. *MIS Q*, 38(1):73–100, March 2014.
- [32] Jean-Baptiste Michel, Yuan Kui Shen, Aviva P. Aiden, Adrian Veres, Matthew K. Gray, The Google Books Team, Joseph P. Pickett, Dale Hoiberg, Dan Clancy,

- Peter Norvig, Jon Orwant, Steven Pinker, Martin A. Nowak, and Erez Lieberman Aiden. Quantitative Analysis of Culture Using Millions of Digitized Books. *Science*, 2010.
- [33] Rada Mihalcea and Vivi Nastase. Word epoch disambiguation: Finding how words change over time. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics, ACL '12*, pages 259–263, 2012.
 - [34] D Gary Miller. *English lexicogenesis*. OUP Oxford, 2014.
 - [35] Sunny Mitra, Ritwik Mitra, Martin Riedl, Chris Biemann, Animesh Mukherjee, and Pawan Goyal. That’s sick dude!: Automatic identification of word sense change across different timescales. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics*, 2014.
 - [36] NSTC. *Nineteenth Century Short Title Catalogue, Series I, Phase 1, 1801-1815*, volume 1. Avero, 1984.
 - [37] Geoff Nunberg. Google books: A metadata train wreck. <http://languageblog.ldc.upenn.edu/nll/?p=1701>, 2009.
 - [38] Terhi Nurmikko-Fuller, Kevin R. Page, Pip Willcox, Jacob Jett, Chris Maden, Timothy Cole, Colleen Fallaw, Megan Senseney, and J. Stephen Downie. Building complex research collections in digital libraries: A survey of ontology implications. In *Proceedings of the 15th ACM/IEEE-CS Joint Conference on Digital Libraries, JCDL '15*, pages 169–172, New York, NY, USA, 2015. ACM.
 - [39] Liudmila Ostroumova Prokhorenkova, Petr Prokhorenkov, Egor Samosvat, and Pavel Serdyukov. Publication date prediction through reverse engineering of the web. In *Proceedings of the Ninth ACM International Conference on Web Search and Data Mining, WSDM '16*, pages 123–132, 2016.
 - [40] Eitan A. Pechenick, Christopher M. Danforth, and Peter Sheridan Dodds. Characterizing the google books corpus: Strong limits to inferences of socio-cultural and linguistic evolution. *PLoS ONE*, 10:e0137041, 2015.
 - [41] Octavian Popescu and Carlo Strapparava. Semeval 2015, task 7: Diachronic text evaluation. In Daniel M. Cer, David Jurgens, Preslav Nakov, and Torsten Zesch, editors, *Proceedings of the 9th International Workshop on Semantic Evaluation (SemEval)*, pages 870–878, 2015.
 - [42] Narayanan Shivakumar and Hector Garcia-Molina. The scam approach to copy detection in digital libraries. Technical report, Stanford InfoLab, 1995.
 - [43] David Smith, Ryan Cordell, and Elizabeth Maddock Dillon. Infectious texts: Modeling text reuse in nineteenth-century newspapers. In *IEEE International Conference on Big Data*, pages 86–94, 2013.
 - [44] David A Smith, Ryan Cordell, Elizabeth Maddock Dillon, Nick Stramp, and John Wilkerson. Detecting and modeling local text reuse. In *Proceedings of the 14th ACM/IEEE-CS Joint Conference on Digital Libraries*, pages 183–192. IEEE Press, 2014.
 - [45] Michael Suarez. Towards a bibliometric analysis of the surviving record, 1701–1800, volume 5 of *The Cambridge History of the Book in Britain*, pages 37–65. Cambridge University Press, 2009.
 - [46] Ted Underwood. Understanding genre in a collection of a million volume. Technical report, University of Illinois, Urbana-Champaign, 2014.
 - [47] Ted Underwood. The life cycles of genres. *Cultural Analytics*, 2016.
 - [48] Matthew Wilkens. The geographic imagination of Civil War-era American fiction. *American Literary History*, 25(4):803–840, 2013.
 - [49] Kyle Williams and C Lee Giles. Near duplicate detection in an academic digital library. In *Proceedings of the 2013 ACM symposium on Document engineering*, pages 91–94. ACM, 2013.
 - [50] Ismet Zeki Yalniz, Ethem F Can, and R Manmatha. Partial duplicate detection for large book collections. In *Proceedings of the 20th ACM international conference on Information and knowledge management*, pages 469–474. ACM, 2011.